



# 쿠버네티스 프로젝트



---

이정일  
서재권  
김민호  
임재근

# 01 [ETCD 백업]

1. <https://127.0.0.1:2379>에서 실행 중인 etcd의 snapshot을 생성하고 snapshot을 etcd-snapshot.db에 저장

후 스냅샷 복원

2. etcdctl을 사용하여 서버에 연결하기 위해 다음 TLS 인증서/키가 제공됩니다.

- CA certificate: /etc/kubernetes/pki/etcd/ca.crt
- Client certificate: /etc/kubernetes/pki/etcd/server.crt
- Client key: /etc/kubernetes/pki/etcd/server.key

1. kubernetes.io/docs에서 etcd backup 검색 후 명령어 찾기

2. 루트 계정으로 전환

```
# sudo -i
```

3. ETCD 클라이언트 설치

```
# apt install etcd-client
```

4. 명령어 수정

```
#ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
--cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kubernetes/pki/etcd/server.crt --key=/etc/kubernetes/pki/etcd/server.key \
snapshot save etcd-snapshot.db
```

5. 위 명령어 적용(백업)

```
root@master:~# ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
--cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kubernetes/pki/etcd/server.crt --key=/etc/kubernetes/pki/etcd/server.key \
snapshot save etcd-snapshot.db
{"level": "info", "ts": 1754896487.0092126, "caller": "snapshot/v3_snapshot.go:119", "msg": "created temporary db file", "path": "etcd-snapshot.db.part"}
{"level": "info", "ts": 2025-08-11T07:14:47.9172082, "caller": "clientv3/maintenance.go:212", "msg": "opened snapshot stream: downloading"}
{"level": "info", "ts": 1754896487.017287, "caller": "snapshot/v3_snapshot.go:127", "msg": "fetching snapshot", "endpoint": "https://127.0.0.1:2379"}
{"level": "info", "ts": 2025-08-11T07:14:48.0140382, "caller": "clientv3/maintenance.go:220", "msg": "completed snapshot read: closing"}
{"level": "info", "ts": 1754896488.0295432, "caller": "snapshot/v3_snapshot.go:142", "msg": "fetched snapshot", "endpoint": "https://127.0.0.1:2379", "size": "7.9 MB", "took": 0.119854689}
{"level": "info", "ts": 1754896488.029675, "caller": "snapshot/v3_snapshot.go:152", "msg": "saved", "path": "etcd-snapshot.db"}
Snapshot saved at etcd-snapshot.db
```

6. 복원

```
root@master:~# ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
--cacert=/etc/kubernetes/pki/etcd/ca.crt --cert=/etc/kubernetes/pki/etcd/server.crt --key=/etc/kubernetes/pki/etcd/server.key \
snapshot restore etcd-snapshot.db
{"level": "info", "ts": 1754896567.1978636, "caller": "snapshot/v3_snapshot.go:386", "msg": "restoring snapshot", "path": "etcd-snapshot.db", "wal-dir": "default.etcd/member/wal", "data-dir": "default.etcd", "snap-dir": "default.etcd/member/snap"}
{"level": "info", "ts": 1754896567.244799, "caller": "mvcc/kvstore.go:388", "msg": "restored last compact revision", "meta-bucket-name": "meta", "meta-bucket-name-key": "finishedCompactRev", "restored-compact-revision": 1286944}
{"level": "info", "ts": 1754896567.2539272, "caller": "membership/cluster.go:392", "msg": "added member", "cluster-id": "cdf818194e3a8c32", "local-member-id": "0", "added-peer-id": "8e9e05c52164694d", "added-peer-peer-urls": ["http://localhost:2380"]}
{"level": "info", "ts": 1754896567.2591524, "caller": "snapshot/v3_snapshot.go:326", "msg": "restored snapshot", "path": "etcd-snapshot.db", "wal-dir": "default.etcd/member/wal", "data-dir": "default.etcd", "snap-dir": "default.etcd/member/snap"}
```

7. tree 설치

```
# apt install tree
```

8. 복원 여부 확인

```
root@master:~#
root@master:~# tree default.etcd/member
default.etcd/member
└── snap
    └── 0000000000000001-0000000000000001.snap
└── db
└── wal
    └── 0000000000000000-0000000000000000.wal
3 directories, 3 files
root@master:~#
```

# 02 [Cluster Upgrade]

1. 마스터 노드의 모든 Kubernetes control plane 및 node 구성 요소를 버전 1.28.8-1.1 버전으로 업그레이드 합니다.
2. master 노드를 업그레이드하기 전에 drain 하고 업그레이드 후에 uncordon해야 합니다.
3. "주의사항" 반드시 Master Node에서 root권한을 가지고 작업을 실행해야 한다.

1. kubectl get node 버전 확인

```
ubuntu@master:~$ kubectl get node
NAME      STATUS   ROLES      AGE      VERSION
master    Ready    control-plane   5d5h    v1.30.14
worker1   Ready    <none>     5d5h    v1.30.14
worker2   Ready    <none>     5d5h    v1.30.14
```

2. kubeadm upgrade검색

3. 업그레이드 버전 결정

```
# sudo apt update
# sudo apt-cache madison kubeadm
```

4. kubeadm 업그레이드 호출

```
# sudo apt-mark unhold kubeadm
```

```
ubuntu@master:~$ sudo apt-mark unhold kubeadm
Canceled hold on kubeadm.
```

```
# sudo apt-get update && sudo apt-get install -y kubeadm='1.30.14-*'
# sudo apt-mark hold kubeadm
```

```
ubuntu@master:~$ sudo apt-mark hold kubeadm
kubeadm set on hold.
```

# kubectl get node

```
ubuntu@master:~$ kubectl get node
NAME      STATUS   ROLES      AGE      VERSION
master    Ready    control-plane   5d6h    v1.30.14
worker1   Ready    <none>     5d6h    v1.30.14
worker2   Ready    <none>     5d6h    v1.30.14
```

5. 업그레이드 버전 선택 후 실행

```
# sudo kubeadm upgrade plan
# sudo kubeadm upgrade apply v1.30.14 --v5
```

6. 노드를 예약 불가능으로 표시하고 유지 관리 준비

```
# kubectl drain master --ignore-daemonsets
```

```
ubuntu@master:~$ kubectl drain master --ignore-daemonsets
node/master cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/calico-node-r5qwx, kube-system/kube-pro
xy-nb66m
evicting pod kube-system/coredns-55cb58b774-xws4m
evicting pod kube-system/coredns-55cb58b774-2mb16
pod/coredns-55cb58b774-2mb16 evicted
pod/coredns-55cb58b774-xws4m evicted
node/master drained
```

7. kubelet 및 kubectl 업그레이드

```
# sudo apt-mark unhold kubelet kubectl
# sudo apt-get update && sudo apt-get install -y kubelet='1.30.14-*' kubectl='1.30.14-*'
# sudo apt-mark hold kubelet kubectl
```

```
ubuntu@master:~$ sudo apt-mark hold kubelet kubectl
kubelet set on hold.
kubectl set on hold.
```

8. kubelet 다시 시작

```
# sudo systemctl daemon-reload
```

```
# sudo systemctl restart kubelet
```

9. 노드 차단 해제

3-15. kubectl uncordon master

```
ubuntu@master:~$ kubectl uncordon master
node/master uncordoned
```

10. kubectl get node 버전 재확인

```
# kubectl get no
```

```
ubuntu@master:~$ kubectl get no
NAME      STATUS    ROLES      AGE      VERSION
master    Ready     control-plane   5d6h    v1.30.14
worker1   Ready     <none>     5d6h    v1.30.14
worker2   Ready     <none>     5d6h    v1.30.14
ubuntu@master:~$
```

# 03 [Service Account, Role, RoleBinding 생성]

애플리케이션 운영중 특정 namespace의 Pod들을 모니터할수 있는 서비스가 요청되었습니다.  
api-access 네임스페이스의 모든 pod를 view할 수 있도록 다음의 작업을 진행하시오.

1. api-access라는 새로운 namespace에 pod-viewer라는 이름의 Service Account를 만듭니다.
2. podreader-role이라는 이름의 Role과 podreader-rolebinding이라는 이름의 RoleBinding을 만듭니다.
3. 앞서 생성한 ServiceAccount를 API resource Pod에 대하여 watch, list, get을 허용하도록 매핑하시오.

1. api-access라는 새로운 namespace 생성

```
# kubectl create ns api-access
```

```
ubuntu@master:~/k8s$ kubectl create namespace api-access
namespace/api-access created
ubuntu@master:~/k8s$ █
```

```
ubuntu@master:~/k8s$ kubectl get namespace
NAME          STATUS   AGE
api-access    Active   20s
default       Active   2d3h
kube-node-lease Active   2d3h
kube-public   Active   2d3h
kube-system   Active   2d3h
ubuntu@master:~/k8s$ █
```

2. pod-viewer라는 이름의 Service Account 생성 및 확인

```
# kubectl create serviceaccount pod-viewer -n api-access
# kubectl get serviceaccount -n api-access
```

```
ubuntu@master:~/k8s$ k create serviceaccount pod-viewer -n api-access
serviceaccount/pod-viewer created
ubuntu@master:~/k8s$ k get serviceaccount -n api-access
NAME      SECRETS   AGE
default   0          10m
pod-viewer 0          8m
ubuntu@master:~/k8s$ █
```

3. watch,list,get 허용하도록 Podreader-role라는 이름의 Role생성

```
# kubectl create role podreader-role -n api-access --resource=pod --verb=watch,list,get
# kubectl describe role -namespace api-access
```

```
ubuntu@master:~/k8s$ kubectl create role podreader-role -n api-access --resource=pod --verb=
watch,list,get
role.rbac.authorization.k8s.io/podreader-role created
ubuntu@master:~/k8s$ █
```

```
ubuntu@master:~/k8s$ kubectl get role -n api-access
NAME      CREATED AT
podreader-role  2025-08-08T04:56:00Z
ubuntu@master:~/k8s$ █
```

4. podreader-rolebinding라는 이름의 RoleBinding 생성

```
# kubectl create rolebinding podreader-rolebinding --serviceaccount=api-account:pod-viewer --role=podreader-
role -n api-access
```

```
# kubectl get rolebinding -namespace api-access
```

```
ubuntu@master:~/k8s$ kubectl create rolebinding podreader-rolebinding --serviceaccount=api-a
ccount:pod-viewer --role=podreader-role -n api-access
rolebinding.rbac.authorization.k8s.io/podreader-rolebinding created
ubuntu@master:~/k8s$ kubectl get rolebinding -n api-access
NAME      ROLE      AGE
podreader-rolebinding  Role/podreader-role  53s
ubuntu@master:~/k8s$ █
```

## 04 [Service Account, ClusterRole, ClusterRoleBinding 생성 생성]

1. 애플리케이션 배포를 위해 새로운 ClusterRole을 생성하고 특정 namespace의 ServiceAccount를 바인드 하시오.
  2. 다음의 resource type에서만 Create가 허용된 ClusterRole deployment-clusterrole을 생성합니다.
  3. Resource Type: Deployment StatefulSet DaemonSet
  4. 미리 생성된 namespace api-access 에 cicd-token이라는 새로운 ServiceAccount를 만듭니다.
  5. ClusterRole deployment-clusterrole을 namespace api-access 로 제한된 새 ServiceAccount cicd-token에 바인딩하세요

## 1. service account 생성

```
# kubectl create serviceaccount cicd-token -n api-access
ubuntu@master:~/k8s$ kubectl create serviceaccount cicd-token -n api-access
serviceaccount/cicd-token created
ubuntu@master:~/k8s$ kubectl get serviceaccount -n api-access
NAME      SECRETS   AGE
cicd-token  0         13s
default     0         24m
pod-viewer  0         22m
ubuntu@master:~/k8s$
```

## 2. cluster role 생성

### 3. cluster role binding 생성

# 05 [노드 비우기]

1. k8s-worker2 노드를 스케줄링 불가능하게 설정
2. 해당 노드에서 실행 중인 모든 Pod을 다른 node로 reschedule

1. worker2 노드 동작 확인

```
# kubectl get nodes
```

```
ubuntu@master:~/k8s$ kubectl get nodes
NAME      STATUS    ROLES          AGE     VERSION
master    Ready     control-plane  2d3h   v1.30.14
worker1   Ready     <none>        2d3h   v1.30.14
worker2   Ready     <none>        2d3h   v1.30.14
ubuntu@master:~/k8s$ █
```

2. Worker2 노드 드레인

```
# kubectl drain worker2 --ignore-daemonsets
```

```
ubuntu@master:~/k8s$ kubectl drain worker2 --ignore-daemonsets
node/worker2 cordoned
Warning: ignoring DaemonSet-managed Pods: kube-system/calico-node-rjfk, kube-system/kube-pro
xy-g6lcp
node/worker2 drained
ubuntu@master:~/k8s$ █
ubuntu@master:~/k8s$ kubectl get nodes
NAME      STATUS    ROLES          AGE     VERSION
master    Ready     control-plane  2d3h   v1.30.14
worker1   Ready     <none>        2d3h   v1.30.14
worker2   Ready,SchedulingDisabled  <none>        2d3h   v1.30.14
ubuntu@master:~/k8s$ █
```

3. Worker2 노드 리스케줄링

```
# kubectl uncordon worker2
```

```
ubuntu@master:~/k8s$ kubectl uncordon worker2
node/worker2 uncordoned
ubuntu@master:~/k8s$ █
```

4. Worker2 노드 상태 확인

```
# kubectl get nodes
```

```
ubuntu@master:~/k8s$ kubectl get nodes
NAME      STATUS    ROLES          AGE     VERSION
master    Ready     control-plane  2d3h   v1.30.14
worker1   Ready     <none>        2d3h   v1.30.14
worker2   Ready     <none>        2d3h   v1.30.14
ubuntu@master:~/k8s$ █
```

# 06 [PodScheduling]

1. 다음의 조건으로 pod를 생성하세요.

- Name: eshop-store
- Image: nginx
- NodeSelector: disktype=ssd

1. 검색 kubernetes.io/docs에서 nodeSelector 검색 후 명령어 찾기

2. Worker1,2 노드에 disktype=ssd & disktype=hdd 라벨 부여

- kubectl label node worker1 disktype=ssd
- kubectl label node worker2 disktype=hdd

라벨 부여 전

```
ubuntu@master:~/k8s$ k get no --show-labels
NAME    STATUS   ROLES      AGE    VERSION   LABELS
master  Ready    control-plane   2d4h   v1.30.14   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,
kubernetes.io/arch=amd64,kubernetes.io/hostname=master,kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=<none>,node.kubernetes.io/exclude-from-external-load-balancers=
worker1 Ready    <none>     2d3h   v1.30.14   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,
kubernetes.io/arch=amd64,kubernetes.io/hostname=worker1,kubernetes.io/os=linux
worker2 Ready    <none>     2d3h   v1.30.14   beta.kubernetes.io/arch=amd64,beta.kubernetes.io/os=linux,
kubernetes.io/arch=amd64,kubernetes.io/hostname=worker2,kubernetes.io/os=linux
ubuntu@master:~/k8s$
```

3. Worker1,2 노드 라벨 확인

```
#kubectl get no -L disktype
```

```
ubuntu@master:~/k8s$ k get no -L disktype
NAME    STATUS   ROLES      AGE    VERSION   DISKTYPE
master  Ready    control-plane   2d4h   v1.30.14
worker1 Ready    <none>     2d3h   v1.30.14   ssd
worker2 Ready    <none>     2d3h   v1.30.14   hdd
ubuntu@master:~/k8s$
```

4. 1. eshop-store.yaml 파일 생성

```
# kubectl run eshop-store --image=nginx --dry-run=client -o yaml > eshop-store.yaml
```

```
ubuntu@master:~/k8s$ kubectl run eshop-store --image=nginx --dry-run=client -o yaml > eshop-store.yaml
ubuntu@master:~/k8s$ vi eshop-store.yaml
```

5. vi로 eshop-store.yaml 파일 내용 수정(nodeSelector 추가)

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: eshop-store
    name: eshop-store
spec:
  containers:
  - image: nginx
    name: eshop-store
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

→

```
apiVersion: v1
kind: Pod
metadata:
  name: eshop-store
spec:
  containers:
  - image: nginx
    name: eshop-store
    nodeSelector:
      disktype: ssd
```

6. eshop-store.yaml 파일 적용

```
ubuntu@master:~/k8s$ k apply -f eshop-store.yaml
pod/eshop-store created
ubuntu@master:~/k8s$ k get po -o wide
NAME        READY   STATUS    RESTARTS   AGE   IP           NODE   NOMINATED-NODE   READINESS   GATES
eshop-store  1/1     Running   0          11s   172.17.235.177   worker1   <none>           <none>
```

# 07 [환경변수, command, args 적용]

1. 'cka-exam'이라는 namespace를 만들고, 아래와 같은 Pod를 생성

- pod Name: pod-01
- image: busybox
- 환경변수 : CERT = "CKA-cert"
- command: /bin/sh
- args: "-c", "while true; do echo \$(CERT); sleep 10;done"

1. NameSpace 생성

```
# kubectl create ns cka-exam
ubuntu@master:~/k8s$ kubectl create ns cka-exam
namespace/cka-exam created
ubuntu@master:~/k8s$ k get ns
NAME      STATUS   AGE
api-access  Active  79m
cka-exam    Active  5s
default     Active  2d4h
kube-node-lease Active  2d4h
kube-public  Active  2d4h
kube-system  Active  2d4h
ubuntu@master:~/k8s$
```

3. pod-01.yaml 파일 생성

```
# kubectl run pod-01 --image=busybox -n cka-exam --env=CERT="CKA-cert" --dry-run=client -o yaml > pod-01.yaml
```

```
ubuntu@master:~/k8s$ kubectl run pod-01 --image=busybox -n cka-exam --env=CERT="CKA-cert" --dry-run=client -o yaml > pod-01.yaml
ubuntu@master:~/k8s$ vi pod-01.yaml
```

4. pod-01.yaml 파일 수정

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: pod-01
  name: pod-01
  namespace: cka-exam
spec:
  containers:
    - env:
        - name: CERT
          value: CKA-cert
      image: busybox
      name: pod-01
      resources: {}
    dnsPolicy: ClusterFirst
    restartPolicy: Always
status: {}
```

→

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-01
  namespace: cka-exam
spec:
  containers:
    - env:
        - name: CERT
          value: CKA-cert
      image: busybox
      name: pod-01
      command: ["/bin/sh"]
      args: ["-c", "while true; do echo $(CERT); sleep 10;done"]
```

5. pod-01.yaml 파일 적용

```
# kubectl apply -f pod-01.yaml
```

```
ubuntu@master:~/k8s$ kubectl apply -f pod-01.yaml
pod/pod-01 created
```

```
ubuntu@master:~/k8s$ kubectl get po -n cka-exam
NAME      READY  STATUS   RESTARTS   AGE
pod-01   1/1    Running   0          99s
```

```
ubuntu@master:~/k8s$ kubectl logs pod-01 -n cka-exam
CKA-cert
CKA-cert
CKA-cert
```

# 08 [Static Pod 생성]

1. worker1 노드에 nginx-static-pod.yaml라는 이름의 Static Pod를 생성

- pod name: nginx-static-pod
- image: nginx
- port : 80

1. nginx-static-pod.yaml 파일 생성

```
- kubectl run nginx-static-pod --image=nginx --port=80 --dry-run=client -o yaml > nginx-static-pod.yaml
```

2. ssh로 worker1 접속

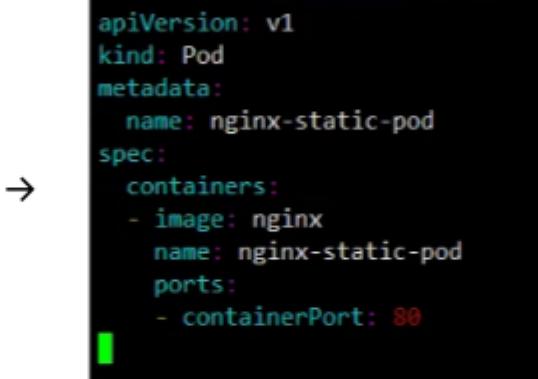
3. worker1의 static위치 확인 및 해당 위치로 이동

```
# cd /var/lib/kubelet  
# cat config.yaml | grep -i static  
# cd /etc/kubernetes/manifests/  
  
ubuntu@worker1:~$ sudo -i  
[sudo] password for ubuntu:  
root@worker1:~# cd /var/lib/kubelet  
root@worker1:/var/lib/kubelet# ls  
checkpoints  cpu_manager_state  kubeadm-flags.env    pki      plugins_registry  pods  
config.yaml  device-plugins     memory_manager_state  plugins  pod-resources  
root@worker1:/var/lib/kubelet# cat config.yaml | grep -i static  
staticPodPath: /etc/kubernetes/manifests  
root@worker1:/var/lib/kubelet#
```

4. nginx-static-pod.yaml 파일 수정 (worker1에 적용)

```
root@worker1:/var/lib/kubelet# cd /etc/kubernetes/manifests  
root@worker1:/etc/kubernetes/manifests# vi nginx-static-pod.yaml  
root@worker1:/etc/kubernetes/manifests#
```

4-1 vi로 worker1에서 내용 수정



The diagram illustrates the modification of a static pod manifest. On the left, the original manifest is shown with a creation timestamp and a null timestamp. An arrow points to the right, where the manifest has been edited. The edited version removes the creation timestamp and changes the timestamp to a current value. The edited manifest is also labeled with a green highlight.

```
apiVersion: v1  
kind: Pod  
metadata:  
  creationTimestamp: null  
  labels:  
    run: nginx-static-pod  
    name: nginx-static-pod  
spec:  
  containers:  
    - image: nginx  
      name: nginx-static-pod  
    ports:  
      - containerPort: 80  
    resources: {}  
  dnsPolicy: ClusterFirst  
  restartPolicy: Always  
status: {}
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx-static-pod  
spec:  
  containers:  
    - image: nginx  
      name: nginx-static-pod  
    ports:  
      - containerPort: 80
```

5. 결과창

```
ubuntu@master:~/k8s$ vi nginx-static-pod.yaml  
ubuntu@master:~/k8s$ k get po  
NAME          READY   STATUS    RESTARTS   AGE  
eshop-store   1/1     Running   0          49m  
nginx-static-pod-worker1 1/1     Running   0          14s  
ubuntu@master:~/k8s$
```

# 09 [로그 확인]

1. Pod "nginx-static-pod-k8s-worker1"의 log를 모니터링하고, 메세지를 포함하는 로그라인을 추출
2. 추출된 결과는 /opt/REPORT/ubuntu/pod-log에 기록

## 1. Pod 확인

```
ubuntu@master:~/k8s$ k get po
NAME                  READY   STATUS    RESTARTS   AGE
eshop-store           1/1     Running   0          49m
nginx-static-pod-worker1 1/1     Running   0          14s
```

## 2. 로그 저장

- kubectl logs nginx-static-pod-worker1 > /opt/REPORT/2023/pod-log

```
ubuntu@master:~$ kubectl logs nginx-static-pod-worker1 > /home/ubuntu/pod-log
ubuntu@master:~$ cat pod-log
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/08/12 06:09:34 [notice] 1#1: using the "epoll" event method
2025/08/12 06:09:34 [notice] 1#1: nginx/1.29.0
2025/08/12 06:09:34 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14+deb12u1)
2025/08/12 06:09:34 [notice] 1#1: OS: Linux 6.8.0-71-generic
2025/08/12 06:09:34 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/08/12 06:09:34 [notice] 1#1: start worker processes
2025/08/12 06:09:34 [notice] 1#1: start worker process 29
2025/08/12 06:09:34 [notice] 1#1: start worker process 30
2025/08/12 06:09:34 [notice] 1#1: start worker process 31
2025/08/12 06:09:34 [notice] 1#1: start worker process 32
```

# 10 [Multi Container Pod 생성]

1. 4개의 컨테이너를 동작시키는 eshop-frontend Pod를 생성하시오.
2. pod image: nginx, redis, memcached, consul

## 1. 앱 생성(eshop-frontend.yaml 파일 생성)

- kubectl run eshop-frontend --image=nginx --dry-run=client -o yaml > multi-pod.yaml

```
ubuntu@master:~/k8s$ kubectl run eshop-frontend --image=nginx --dry-run=client -o yaml > multi-pod.yaml
ubuntu@master:~/k8s$
```

## 2. eshop-frontend.yaml 파일 수정

- vi multi-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: eshop-frontend
    name: eshop-frontend
spec:
  containers:
    - image: nginx
      name: eshop-frontend
      resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```



```
apiVersion: v1
kind: Pod
metadata:
  run: eshop-frontend
  name: eshop-frontend
spec:
  containers:
    - image: nginx
      name: nginx
    - image: redis
      name: redis
    - image: memcached
      name: memcached
    - image: consul
      name: consul
```

## 3. eshop-frontend.yaml 파일 적용

```
ubuntu@master:~/k8s$ k apply -f multi-pod.yaml
pod/eshop-frontend created
ubuntu@master:~/k8s$ k get po
NAME          READY   STATUS        RESTARTS   AGE
eshop-frontend 0/4     ContainerCreating  0          4s
eshop-store    1/1     Running       0          80m
nginx-static-pod-worker1 0/1     Completed    0          30m
ubuntu@master:~/k8s$
```

# 11 [Rolling Update & Roll Back]

1. Deployment를 이용해 nginx 파드를 3개 배포한 다음 컨테이너 이미지 버전을 rolling update하고 update record를 기록

2. 마지막으로 컨테이너 이미지를 previous version으로 roll back

- name: eshop-payment
- Image : nginx
- Image version: 1.16
- update image version: 1.17
- label: app=payment, environment=production

1. eshop-payment.yaml 파일 생성

```
# kubectl create deploy eshop-payment --image=nginx:1.16 --replicas=3 --dry-run=client -o yaml > eshop-payment.yaml
```

```
ubuntu@master:~/k8s$ kubectl create deploy eshop-payment --image=nginx:1.16 --replicas=3 --dry-run=client -o yaml > eshop-payment.yaml
ubuntu@master:~/k8s$
```

2. 1. eshop-payment.yaml 파일 수정

- vi eshop-payment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: eshop-payment
  name: eshop-payment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: eshop-payment
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: eshop-payment
    spec:
      containers:
        - image: nginx:1.16
          name: nginx
          resources: {}
status: {}
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app: payment
    environment: production
  name: eshop-payment
spec:
  replicas: 3
  selector:
    matchLabels:
      app: payment
      environment: production
  template:
    metadata:
      labels:
        app: payment
        environment: production
    spec:
      containers:
        - image: nginx:1.16
          name: nginx
```

3. 1. eshop-payment.yaml 파일 적용(- --record)

```
# kubectl apply -f eshop-payment.yaml - --record
```

```
ubuntu@master:~/k8s$ kubectl apply -f eshop-payment.yaml --record
Flag --record has been deprecated, --record will be removed in the future
error: the path "eshop-payment.yaml" does not exist
```

```
ubuntu@master:~/k8s$ kubectl get deploy,rs,po
NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/eshop-payment         3/3     3           3           102s
replicaset.apps/eshop-payment-6855fb78d6   3       3           3           102s
pod/eshop-payment-6855fb78d6-8ch4s      1/1     Running     0           102s
pod/eshop-payment-6855fb78d6-hlpgd       1/1     Running     0           102s
pod/eshop-payment-6855fb78d6-nlbqw       1/1     Running     0           102s
pod/nginx-static-pod-worker1            1/1     Running     0           4s
```

```
ubuntu@master:~/k8s$ kubectl rollout history deploy eshop-payment
deployment.apps/eshop-payment
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=eshop-payment.yaml --record=true
```

#### 4. 롤링 업데이트

```
# kubectl set image deploy eshop-payment nginx=nginx:1.17 - --record
```

```
ubuntu@master:~/k8s$ kubectl set image deploy eshop-payment nginx=nginx:1.17 --record
Flag --record has been deprecated, --record will be removed in the future
deployment.apps/eshop-payment image updated
ubuntu@master:~/k8s$ kubectl get deploy,rs,po
NAME                      READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/eshop-payment   3/3     3           3          4m32s

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/eshop-payment-6855fb78d6   0         0         0        4m32s
replicaset.apps/eshop-payment-696dd55885   3         3         3        15s

NAME                      READY   STATUS    RESTARTS   AGE
pod/eshop-payment-696dd55885-7jhbc   1/1     Running   0          14s
pod/eshop-payment-696dd55885-fpc8t    1/1     Running   0          15s
pod/eshop-payment-696dd55885-qh2b8    1/1     Running   0          13s
pod/nginx-static-pod-worker1       1/1     Running   0          2m54s
ubuntu@master:~/k8s$ kubectl rollout history deploy eshop-payment
deployment.apps/eshop-payment
REVISION  CHANGE-CAUSE
1          kubectl apply --filename=eshop-payment.yaml --record=true
2          kubectl set image deploy eshop-payment nginx=nginx:1.17 --record=true
```

#### 5. 롤백

```
# kubectl rollout undo deploy eshop-payment
```

```
ubuntu@master:~/k8s$ kubectl rollout undo deploy eshop-payment
deployment.apps/eshop-payment rolled back
ubuntu@master:~/k8s$ kubectl rollout history deploy eshop-payment
deployment.apps/eshop-payment
REVISION  CHANGE-CAUSE
2          kubectl set image deploy eshop-payment nginx=nginx:1.17 --record=true
3          kubectl apply --filename=eshop-payment.yaml --record=true
ubuntu@master:~/k8s$ kubectl describe pod eshop-payment-6855fb78d6-bjjjw | grep -i image
  Image:      nginx:1.16
  Image ID:   docker.io/library/nginx@sha256:d20aa6d1cae56fd17cd458f4807e0de462caf2336
f0b70b5eeb69fcfaaf30dd9c
  Normal  Pulled   2m5s  kubelet      Container image "nginx:1.16" already present on machine
```

# 12 [ClusterIP]

1. 'devops' namespace에서 deployment eshop-order를 다음 조건으로 생성하시오.

- image: nginx, replicas: 2, label: name=order

2. 'eshop-order' deployment의 Service를 만드세요.

- Service Name: eshop-order-svc

- Type: ClusterIP

- Port: 80

1. devops 네임스페이스 생성

```
# kubectl create ns devops
```

```
# kubectl get namespace
```

```
ubuntu@master:~/k8s$ kubectl create namespace devops
namespace/devops created
```

```
ubuntu@master:~/k8s$ kubectl get namespace
```

NAME	STATUS	AGE
api-access	Active	3h13m
cka-exam	Active	114m
default	Active	2d6h
devops	Active	8s
kube-node-lease	Active	2d6h
kube-public	Active	2d6h
kube-system	Active	2d6h

2. eshop-order.yaml 파일 생성

```
# kubectl create deploy eshop-order -n devops - --image=nginx - --replicas=2 - --dry-run=client -o yaml > eshop-order.yaml
```

```
ubuntu@master:~/k8s$ kubectl create deploy eshop-order -n devops --image=nginx --replicas=2
--dry-run=client -o yaml > eshop-order.yaml
```

3. YAML 수정

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: eshop-order
  name: eshop-order
  namespace: devops
spec:
  replicas: 2
  selector:
    matchLabels:
      app: eshop-order
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: eshop-order
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    name: order
  name: eshop-order
  namespace: devops
spec:
  replicas: 2
  selector:
    matchLabels:
      name: order
  strategy: {}
  template:
    metadata:
      labels:
        name: order
    spec:
      containers:
        - image: nginx
          name: nginx
```

#### 4. eshop-order.yaml 파일 적용

```
# kubectl apply -f eshop-order.yaml
```

```
ubuntu@master:~/k8s$ kubectl apply -f eshop-order.yaml
deployment.apps/eshop-order created
ubuntu@master:~/k8s$ kubectl get deploy -n devops
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
eshop-order  2/2     2           2           13s
```

#### 4. 서비스 생성

```
# kubectl expose deploy eshop-order -n devops --name=eshop-order-svc --port=80 --target-port=80
# kubectl get deploy,svc -n devops
```

```
ubuntu@master:~/k8s$ kubectl get deploy,rs,po -n devops
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/eshop-order  2/2     2           2           41s

NAME          DESIRED   CURRENT   READY   AGE
replicaset.apps/eshop-order-d6d584684  2         2         2         41s

NAME          READY   STATUS    RESTARTS   AGE
pod/eshop-order-d6d584684-h8dsj  1/1     Running   0          41s
pod/eshop-order-d6d584684-vh5dd  1/1     Running   0          41s
ubuntu@master:~/k8s$ kubectl expose deploy eshop-order -n devops --name=eshop-order-svc --port=80 --target-port=80
service/eshop-order-svc exposed
ubuntu@master:~/k8s$ kubectl get deploy,svc -n devops
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/eshop-order  2/2     2           2           3m56s

NAME        TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
service/eshop-order-svc  ClusterIP  10.110.199.210  <none>       80/TCP    20s
```

# 13 [NodePort]

1. 'front-end' deployment를 다음 조건으로 생성하시오.
  - image: nginx, replicas: 2, label: run=nginx
2. 'front-end' deployment의 nginx 컨테이너를 expose하는 'front-end-nodesvc'라는 새 service를 만듭니다.
3. Front-end로 동작중인 Pod에는 node의 30200 포트로 접속되어야 합니다.

1. front-end.yaml 파일 생성

```
# kubectl create deploy front-end - -image=nginx - -replicas=2 - -dry-run=client -o yaml > front-end.yaml
```

```
ubuntu@master:~/k8s$ kubectl create deploy front-end --image=nginx --replicas=2 --dry-run=client -o yaml > front-end.yaml
ubuntu@master:~/k8s$
```

2. front-end.yaml 파일 수정

```
# - vi front-end.yaml
# - run=nginx → run: nginx
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: front-end
  name: front-end
spec:
  replicas: 2
  selector:
    matchLabels:
      app: front-end
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: front-end
    spec:
      containers:
        - image: nginx
          name: nginx
          resources: {}
status: {}
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    run: nginx
  name: front-end
spec:
  replicas: 2
  selector:
    matchLabels:
      run: nginx
  template:
    metadata:
      labels:
        run: nginx
    spec:
      containers:
        - image: nginx
          name: nginx
```

3. front-end.yaml 파일 적용

```
# kubectl apply -f front-end.yaml
# kubectl get deploy,rs,po
```

```
ubuntu@master:~/k8s$ vi front-end.yaml
ubuntu@master:~/k8s$ kubectl apply -f front-end.yaml
deployment.apps/front-end created
ubuntu@master:~/k8s$ kubectl get deploy,rs,po
NAME                      READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/front-end   2/2     2           2           10s

NAME                           DESIRED   CURRENT   READY   AGE
replicaset.apps/front-end-7858f98fb4   2         2         2         10s

NAME                     READY   STATUS    RESTARTS   AGE
pod/front-end-7858f98fb4-k8s95   1/1     Running   0          10s
pod/front-end-7858f98fb4-vthxc   1/1     Running   0          10s
ubuntu@master:~/k8s$
```

#### 4. front-end-nodesvc.yaml 파일 생성

```
#kubectl expose deploy front-end - -name=front-end-nodesvc - -port=80 - -target-port=80 - -type=NodePort - -dry-run=client -o yaml > front-end-nodesvc.yaml
```

```
ubuntu@master:~/k8s$ kubectl expose deploy front-end --name=front-end-nodesvc --port=80 --target-port=80 --type=NodePort --dry-run=client -o yaml > front-end-nodesvc.yaml
ubuntu@master:~/k8s$
```

#### 5.front-end-nodesvc.yaml 파일수정

```
# vi front-end-nodesvc.yaml
```

```
apiVersion: v1
kind: Service
metadata:
  creationTimestamp: null
  labels:
    run: nginx
  name: front-end-nodesvc
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
  selector:
    run: nginx
  type: NodePort
status:
  loadBalancer: {}
```

→

```
apiVersion: v1
kind: Service
metadata:
  labels:
    run: nginx
  name: front-end-nodesvc
spec:
  ports:
    - port: 80
      protocol: TCP
      targetPort: 80
      nodePort: 30200
  selector:
    run: nginx
  type: NodePort
```

~

#### 6. front-end-nodesvc.yaml 파일 적용

```
# kubectl apply -f front-end-nodesvc.yaml
```

```
ubuntu@master:~/k8s$ kubectl apply -f front-end-nodesvc.yaml
service/front-end-nodesvc created
ubuntu@master:~/k8s$ kubectl get svc
NAME         TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)      AGE
front-end-nodesvc   NodePort   10.110.167.65  <none>        80:30200/TCP  10s
kubernetes     ClusterIP  10.96.0.1    <none>        443/TCP      16m
ubuntu@master:~/k8s$
```

# 14 [Network Policy]

1. customera, customerb를 생성한 후, 각각 PARTITION=customera, PARTITION=customerb를 라벨링 하시오.

2. default namespace에 다음과 같은 pod를 생성하세요.

```
- name: poc  
- image: nginx  
- port: 80  
- label: app=poc  
- "partition=customera"를 사용하는 namespace에서만 poc의
```

3. 80포트로 연결할 수 있도록 default namespace에 'allow-web-from-customera'라는 network Policy를 설정하세요.

보안 정책상 다른 namespace의 접근은 제한합니다.

1. customera, customerb 네임스페이스 생성 및 확인

```
# kubectl create ns customera  
# kubectl create ns customerb  
# kubectl get ns customera customerb  
ubuntu@master:~/k8s$ kubectl create ns customera  
namespace/customera created  
ubuntu@master:~/k8s$ kubectl create ns customerb  
namespace/customerb created  
ubuntu@master:~/k8s$ kubectl get ns customera customerb  
NAME      STATUS   AGE  
customera  Active   33s  
customerb  Active   28s  
ubuntu@master:~/k8s$
```

2. customera, customerb 네임스페이스 라벨링

```
# kubectl label ns customera PARTITION=customera  
# kubectl label ns customerb PARTITION=customerb  
# kubectl get ns --show-labels  
ubuntu@master:~/k8s$ kubectl label ns customera PARTITION=customera  
namespace/customera labeled  
ubuntu@master:~/k8s$ kubectl label ns customerb PARTITION=customerb  
namespace/customerb labeled  
ubuntu@master:~/k8s$ kubectl get ns --show-labels  
NAME      STATUS   AGE   LABELS  
customera  Active   2m52s  PARTITION=customera,kubernetes.io/metadata.name=customera  
customerb  Active   2m47s  PARTITION=customerb,kubernetes.io/metadata.name=customerb  
default    Active   5d1h   kubernetes.io/metadata.name=default  
kube-node-lease  Active   5d1h   kubernetes.io/metadata.name=kube-node-lease  
kube-public   Active   5d1h   kubernetes.io/metadata.name=kube-public  
kube-system   Active   5d1h   kubernetes.io/metadata.name=kube-system  
ubuntu@master:~/k8s$
```

### 3. 파드 생성

```
# kubectl run poc - --image=nginx - --port=80 - --labels=app=poc
```

```
# kubectl get po poc - --show-labels
```

```
ubuntu@master:~/k8s$ kubectl run poc --image=nginx --port=80 --labels=app=poc
pod/poc created
ubuntu@master:~/k8s$ kubectl get po poc --show-labels
NAME    READY   STATUS    RESTARTS   AGE   LABELS
poc    1/1     Running   0          11s   app=poc
ubuntu@master:~/k8s$
```

### 4. netpol.yaml 파일 생성

```
# 검색 : networkpolicy(쿠버네티스 사이트)
```

### 5. netpol.yaml 파일 수정

```
# vi netpol.yaml
```

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
        - except:
            - cidr: 172.17.1.0/24
    - namespaceSelector:
        matchLabels:
          project: myproject
    - podSelector:
        matchLabels:
          role: frontend
  ports:
    - protocol: TCP
      port: 6379
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
  ports:
    - protocol: TCP
      port: 5978
```



```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-web-from-customer-a
  namespace: default
spec:
  podSelector:
    matchLabels:
      app: poc
  policyTypes:
    - Ingress
  ingress:
    - from:
        - namespaceSelector:
            matchLabels:
              partition: customer-a
  ports:
    - protocol: TCP
      port: 80
```

### 6. netpol.yaml 파일 적용

```
# kubectl apply -f netpol.yaml
```

```
# kubectl get netpol
```

```
ubuntu@master:~/k8s$ kubectl apply -f netpol.yaml
networkpolicy.networking.k8s.io/allow-web-from-customer-a created
ubuntu@master:~/k8s$ kubectl get netpol
NAME                POD-SELECTOR   AGE
allow-web-from-customer-a   app=poc       8s
ubuntu@master:~/k8s$
```

# 15 [Ingress]

1. Create a new nginx Ingress resource as follows:

- Name: ping
- Namespace: ing-internal
- Exposing service hi on path /hi using service port 5678

1. Ing-internal 네임스페이스 생성 및 확인

```
# kubectl create ns ing-internal  
# kubectl get ns
```

```
ubuntu@master:~/k8s$ kubectl create ns ing-internal  
namespace/ing-internal created  
ubuntu@master:~/k8s$ kubectl get ns  
NAME STATUS AGE  
customer-a Active 83m  
customer-b Active 83m  
default Active 5d3h  
ing-internal Active 10s  
kube-node-lease Active 5d3h  
kube-public Active 5d3h  
kube-system Active 5d3h  
ubuntu@master:~/k8s$ █
```

2. ingress.yaml 파일 생성

```
# 검색 : ingress(쿠버네티스 사이트)
```

3. ingress.yaml 파일 수정

```
# vi ingress.yaml
```

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: minimal-ingress  
  annotations:  
    nginx.ingress.kubernetes.io/rewrite-target: /  
spec:  
  ingressClassName: nginx-example  
  rules:  
  - http:  
    paths:  
    - path: /testpath  
      pathType: Prefix  
      backend:  
        service:  
          name: test  
          port:  
            number: 80  
█
```



```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: ping  
  namespace: ing-internal  
spec:  
  ingressClassName: nginx-example  
  rules:  
  - http:  
    paths:  
    - path: /hi  
      pathType: Prefix  
      backend:  
        service:  
          name: hi  
          port:  
            number: 5678
```

4. ingress.yaml 파일 적용

```
# kubectl apply -f ingress.yaml
```

```
# kubectl get ing -n ing-internal
```

```
ubuntu@master:~/k8s$ kubectl apply -f ingress.yaml  
ingress.networking.k8s.io/ping created  
ubuntu@master:~/k8s$ kubectl get ing -n ing-internal  
NAME CLASS HOSTS ADDRESS PORTS AGE  
ping nginx-example * 80 37s  
ubuntu@master:~/k8s$ █
```

# 16 [Service and DNS Lookup]

1. image nginx를 사용하는 resolver pod를 생성하고 resolver-service라는 service를 구성

2. 클러스터 내에서 service와 pod 이름을 조회할 수 있는지 테스트

- dns 조회에 사용하는 pod 이미지는 busybox:1.28이고, service와 pod 이름 조회는 nslookup을 사용
- service 조회 결과는 /home/ubuntu/nginx.svc에 pod name 조회 결과는 /home/ubuntu/nginx.pod 파일에 기록

1. resolver pod 생성 및 확인

```
# kubectl run resolver --image=nginx --port=80  
# kubectl get po resolver
```

```
ubuntu@master:~/k8s$ kubectl run resolver --image=nginx --port=80  
pod/resolver created  
ubuntu@master:~/k8s$ kubectl get po resolver  
NAME      READY   STATUS    RESTARTS   AGE  
resolver   1/1     Running   0          9s  
ubuntu@master:~/k8s$
```

2. 서비스 생성 및 확인

```
# kubectl expose pod resolver --name=resolver-service --port=80  
# kubectl get svc resolver-service
```

```
ubuntu@master:~/k8s$ kubectl expose pod resolver --name=resolver-service --port=80  
service/resolver-service exposed  
ubuntu@master:~/k8s$ kubectl get svc  
NAME        TYPE      CLUSTER-IP    EXTERNAL-IP   PORT(S)      AGE  
front-end-nodesvc  NodePort  10.110.167.65 <none>       80:30200/TCP  108m  
kubernetes   ClusterIP  10.96.0.1    <none>       443/TCP     125m  
resolver-service  ClusterIP  10.102.45.25 <none>       80/TCP      9s  
ubuntu@master:~/k8s$
```

3. 서비스 DNS 조회 및 저장

```
# kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10.102.45.25
```

```
ubuntu@master:~/k8s$ kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10.102.45.25  
5.25  
Server: 10.96.0.10  
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local  
  
Name: 10.102.45.25  
Address 1: 10.102.45.25 resolver-service.default.svc.cluster.local  
pod "testpod" deleted
```

```
# mkdir cka
```

```
# kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10.102.45.25 > nginx.svc  
# ls
```

```
ubuntu@master:~/k8s$ mkdir cka  
ubuntu@master:~/k8s$ kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10.102.45.25 > /home/ubuntu/cka/nginx.svc  
-bash: /home/ubuntu/cka/nginx.svc: No such file or directory  
ubuntu@master:~/k8s$ ls  
annotation.yaml  development-label01.yaml  front-end-nodesvc.yaml  loadbalancer.yaml  service.yaml  
cka              ex.yaml                  front-end.yaml        netpol.yaml  
deployment-nginx.yaml  front                  ingress.yaml       nodeport.yaml  
ubuntu@master:~/k8s$
```

```
# cat nginx.svc
```

```
ubuntu@master:~/k8s$ cat nginx.svc
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10.102.45.25
Address 1: 10.102.45.25 resolver-service.default.svc.cluster.local
pod "testpod" deleted
ubuntu@master:~/k8s$
```

#### 4. 파드 DNS 조회 및 저장

```
# kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10-101-78-250.default.pod.cluster.local
# kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10-101-78-250.default.pod.cluster.local > nginx.pod
# cat nginx.pod
```

```
ubuntu@master:~/k8s$ kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10-101-78-250.default.pod.cluster.local
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10-101-78-250.default.pod.cluster.local
Address 1: 10.101.78.250
pod "testpod" deleted
ubuntu@master:~/k8s$ kubectl run testpod --image=busybox:1.28 -it --restart=Never --rm --nslookup 10-101-78-250.default.pod.cluster.local > nginx.pod
ubuntu@master:~/k8s$ cat nginx.pod
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: 10-101-78-250.default.pod.cluster.local
Address 1: 10.101.78.250
pod "testpod" deleted
```

# 17 [emptyDir Volume]

1. 다음 조건에 맞춰서 nginx 웹서버 pod가 생성한 로그파일을 받아서 STDOUT으로 출력하는 busybox 컨테이너를 운영

Pod Name: weblog

Web container:

- Image: nginx:1.17
- Volume mount : /var/log/nginx

- Readwrite

Log container:

- Image: busybox
- args: /bin/sh, -c, "tail -n+1 -f /data/access.log"
- Volume mount : /data
- readonly

1. weblog.yaml 생성

```
# kubectl run weblog --image=nginx:1.17 --dry-run=client -o yaml > weblog.yaml
```

2. weblog.yaml 수정

```
# vi weblog.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: weblog
    name: weblog
spec:
  containers:
    - image: nginx:1.17
      name: weblog
      resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```



```
apiVersion: v1
kind: Pod
metadata:
  name: weblog
spec:
  containers:
    - image: nginx:1.17
      name: web
      volumeMounts:
        - mountPath: /var/log/nginx
          name: weblog
    - image: busybox
      args: [/bin/sh, -c, "tail -n+1 -f /data/access.log"]
      name: log
      volumeMounts:
        - mountPath: /data
          name: weblog
          readOnly: true
  volumes:
    - name: weblog
      emptyDir: {}
```

3. weblog yaml 파일 적용

```
# kubectl apply -f weblog.yaml
```

```
ubuntu@master:~/k8s$ kubectl apply -f weblog.yaml
pod/weblog created
ubuntu@master:~/k8s$ kubectl get pod weblog
NAME      READY   STATUS    RESTARTS   AGE
weblog   2/2     Running   0          7s
Containers:
  weblog:
    Container ID:  containerd://7f9c0ab4bea18d17625925f576a02c58d320f8f78de49b5c8e4cb1d2c177
    d16f
    Image:          nginx:1.17
    Image ID:       docker.io/library/nginx@sha256:6fff55753e3b34e36e24e37039ee9eae1fe38a6420
    d8e16ef37c92d1eb26699
    Port:           <none>
    Host Port:      <none>
    State:          Running
    Started:        Mon, 11 Aug 2025 06:07:55 +0000
    Ready:          True
    Restart Count:  0
    Environment:    <none>
    Mounts:
      /var/log/nginx from weblog (rw)
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-5rlc5 (ro)
  log:
    Container ID:  containerd://1c0e459c4ad3bbaf1b49c22b54a5a20b08585b7bd9c4249ddfb0a4efce4ea9
    08e
    Image:          nginx:1.17
    Image ID:       docker.io/library/nginx@sha256:6fff55753e3b34e36e24e37039ee9eae1fe38a6420
    d8e16ef37c92d1eb26699
    Port:           <none>
    Host Port:      <none>
    Args:
      /bin/sh
      -c
      tail -n+1 -f /data/access.log
    State:          Running
    Started:        Mon, 11 Aug 2025 06:07:55 +0000
    Ready:          True
    Restart Count:  0
    Environment:    <none>
    Mounts:
      /data from weblog (ro)
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-5rlc5 (ro)
Conditions:
  Type            Status
  PodReadyToStartContainers  True
```

# 18 [HostPath Volume]

1. /data/cka/fluentd.yaml 파일을 만들어 새로운 Pod 생성하세요

- 신규생성 Pod Name: fluentd, image: fluentd, namespace: default)

2. 위 조건을 참고하여 다음 조건에 맞게 볼륨마운트를 설정하시오.

3. Worker node의 도커 컨테이너 디렉토리 : /var/lib/docker/containers 동일 디렉토리로 pod에 마운트 하시오.
4. Worker node의 /var/log 디렉토리를 fluentd Pod에 동일이름의 디렉토리 마운트하시오.

1. 검색 : hostpath 검색( (쿠버네티스 사이트)

2. cka 폴더 생성 및 이동

```
# mkdir cka  
# cd cka
```

```
ubuntu@master:~$ mkdir cka  
ubuntu@master:~$ cd cka  
ubuntu@master:~/cka$ █
```

3. fluentd pod 파드 생성

```
# kubectl run fluentd - --image=fluentd - --port=80
```

```
ubuntu@master:~/k8s$ kubectl run fluentd --image=fluentd --port=80  
pod/fluentd created  
ubuntu@master:~/k8s$ █
```

4. fluentd yaml파일 생성

```
# kubectl get po fluentd -o yaml > fluentd.yaml
```

```
ubuntu@master:~/cka$ kubectl get po fluentd -o yaml > fluentd.yaml  
ubuntu@master:~/cka$ █
```

5. fluentd yaml 파일 수정

```
# vi fluentd.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  annotations:  
    csi.projectcalico.org/containerID: 5D0089A09A8F12861792295e291ce73109e4004554517a54f9e1e29d9fe3d  
    csi.projectcalico.org/podIP: 172.16.235.54/32  
    csi.projectcalico.org/podIPs: 172.16.235.141/32  
  creationTimestamp: "2025-06-11T06:17:00Z"  
  labels:  
    run: fluentd  
  name: fluentd  
  namespace: default  
  resourceVersion: "44337"  
  uid: 0179038E-9c6f-45de-b7fc-5d57d8d4d85  
spec:  
  containers:  
    - image: fluentd  
      imagePullPolicy: Always  
      name: fluentd  
      ports:  
        - containerPort: 80  
          protocol: TCP  
      resources: {}  
      terminationMessagePath: /dev/termination-log  
      terminationMessagePolicy: Final  
      volumeMounts:  
        - mountPath: /var/run/secrets/kubernetes.io/serviceaccount  
          name: kube-api-access-egxz  
          readOnly: true  
      dnsPolicy: ClusterFirst  
      enableServiceLinks: true  
      sidecars: []  
      terminationMessagePath: /dev/termination-log  
      terminationMessagePolicy: Final  
      priority: 0  
      preemptionPolicy: PreemptLowerPriority  
      priorityClassName: null  
      restartPolicy: Always  
      schedulerName: default-scheduler  
      securityContext: {}  
  "fluentd.yaml" 11L, 3029
```



```
apiVersion: v1  
kind: Pod  
metadata:  
  name: fluentd  
spec:  
  containers:  
    - image: fluentd  
      name: fluentd  
      ports:  
        - containerPort: 80  
          protocol: TCP  
      volumeMounts:  
        - mountPath: /var/lib/docker/containers  
          name: containersdir  
        - mountPath: /var/log  
          name: logdir  
      volumes:  
        - name: containersdir  
          hostPath:  
            path: /var/lib/docker/containers  
        - name: logdir  
          hostPath:  
            path: /var/log
```

## 5. fluentd yaml 파일 적용, 파드 삭제

```
# kubectl delete po fluentd --force
```

```
ubuntu@master:~/cka$ kubectl delete po fluentd --force
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "fluentd" force deleted
```

## 6. fluentd yaml 파일 적용

```
# kubectl apply -f fluentd.yaml
```

```
# kubectl describe pod fluentd
```

```
ubuntu@master:~/cka$ kubectl apply -f fluentd.yaml
pod/fluentd created
ubuntu@master:~/cka$ kubectl get pod
NAME      READY   STATUS    RESTARTS   AGE
fluentd   1/1     Running   0          6s
weblog    2/2     Running   0          37m
```

# 19 [Persistent Volume]

1. pv001라는 이름으로 size 1Gi, access mode ReadWriteMany를 사용하여 persistent volume을 생성합니다.

2. volume type은 hostPath이고 위치는 /tmp/app-config입니다.

1. hostpath 검색

```
# kind : Persistent
```

```
2. pv001.yaml 파일 복사 붙여넣기
```

```
# vi pv001.yaml
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv0003
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteOnce
~
```

3. hostpath 검색 후 복사

```
# hostPath:
```

```
  path: /any/path/it/will/be/replaced
```

4.pv001.yaml 파일 수정

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv001
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  hostPath:
    path: /tmp/app-config
~
```

5. pv001.yaml 파일 적용

```
# kubectl apply -f pv001.yaml
```

```
# kubectl get pv
```

```
ubuntu@master:~$ kubectl apply -f pv001.yaml
persistentvolume/pv001 created
ubuntu@master:~$ kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS     CLAIM   STORAGECLASS   VOLUME ATTRIBUTESCLASS
pv001      1Gi        Rwo          Retain           Available   <none>  <none>
9s
ubuntu@master:~$
```



대우직업능력개발원

Deawon Development Institute for Vocational ability



감사합니다

APEX